

```
In [105... #Importing Libraries needed in the project
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

```
In [106... #Reading the dataset from the excel and creating a dataframe
# Printing First 5 rows
df_housing = pd.read_excel('1553768847_housing.xlsx')
df_housing.head()
```

```
Out[106... longitude latitude housing_median_age total_rooms total_bedrooms population households med
```

0	-122.23	37.88	41	880	129.0	322	126
1	-122.22	37.86	21	7099	1106.0	2401	1138
2	-122.24	37.85	52	1467	190.0	496	177
3	-122.25	37.85	52	1274	235.0	558	219
4	-122.25	37.85	52	1627	280.0	565	259

```
In [107... # Checking the mean/ counts and std in the dataframe
df_housing.describe()
```

```
Out[107... longitude latitude housing_median_age total_rooms total_bedrooms population
```

count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744	
std	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.462122	
min	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000	
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000	
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000	
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000	
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	

```
In [108... # Checking the data for null columns and datatypes to see if any non numrical column ex
df_housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype
---  -
0   longitude            20640 non-null  float64
```

```

1  latitude                20640 non-null float64
2  housing_median_age      20640 non-null int64
3  total_rooms             20640 non-null int64
4  total_bedrooms          20433 non-null float64
5  population              20640 non-null int64
6  households              20640 non-null int64
7  median_income           20640 non-null float64
8  ocean_proximity         20640 non-null object
9  median_house_value       20640 non-null int64
dtypes: float64(4), int64(5), object(1)
memory usage: 1.6+ MB

```

```

In [109... #count of Null values
df_housing.isnull().sum()

```

```

Out[109... longitude                0
latitude                0
housing_median_age      0
total_rooms             0
total_bedrooms          207
population              0
households              0
median_income           0
ocean_proximity         0
median_house_value      0
dtype: int64

```

```

In [110... #Fill the missing values with the mean of the respective column.
df_housing.fillna(df_housing['total_bedrooms'].mean(), inplace=True)

```

```

In [111... df_housing.isnull().sum()

```

```

Out[111... longitude                0
latitude                0
housing_median_age      0
total_rooms             0
total_bedrooms          0
population              0
households              0
median_income           0
ocean_proximity         0
median_house_value      0
dtype: int64

```

```

In [112... #Convert categorical column in the dataset to numerical data using one hot encoder
df_housing_final = pd.get_dummies(data= df_housing, columns=['ocean_proximity'])
df_housing_final.shape

```

```

Out[112... (20640, 14)

```

```

In [113... #Extract input (X) and output (Y) data from the dataset.
x_input = df_housing_final.drop(columns=['median_house_value'])
y_output = df_housing_final['median_house_value']
print(x_input.shape,y_output.shape)

```

```

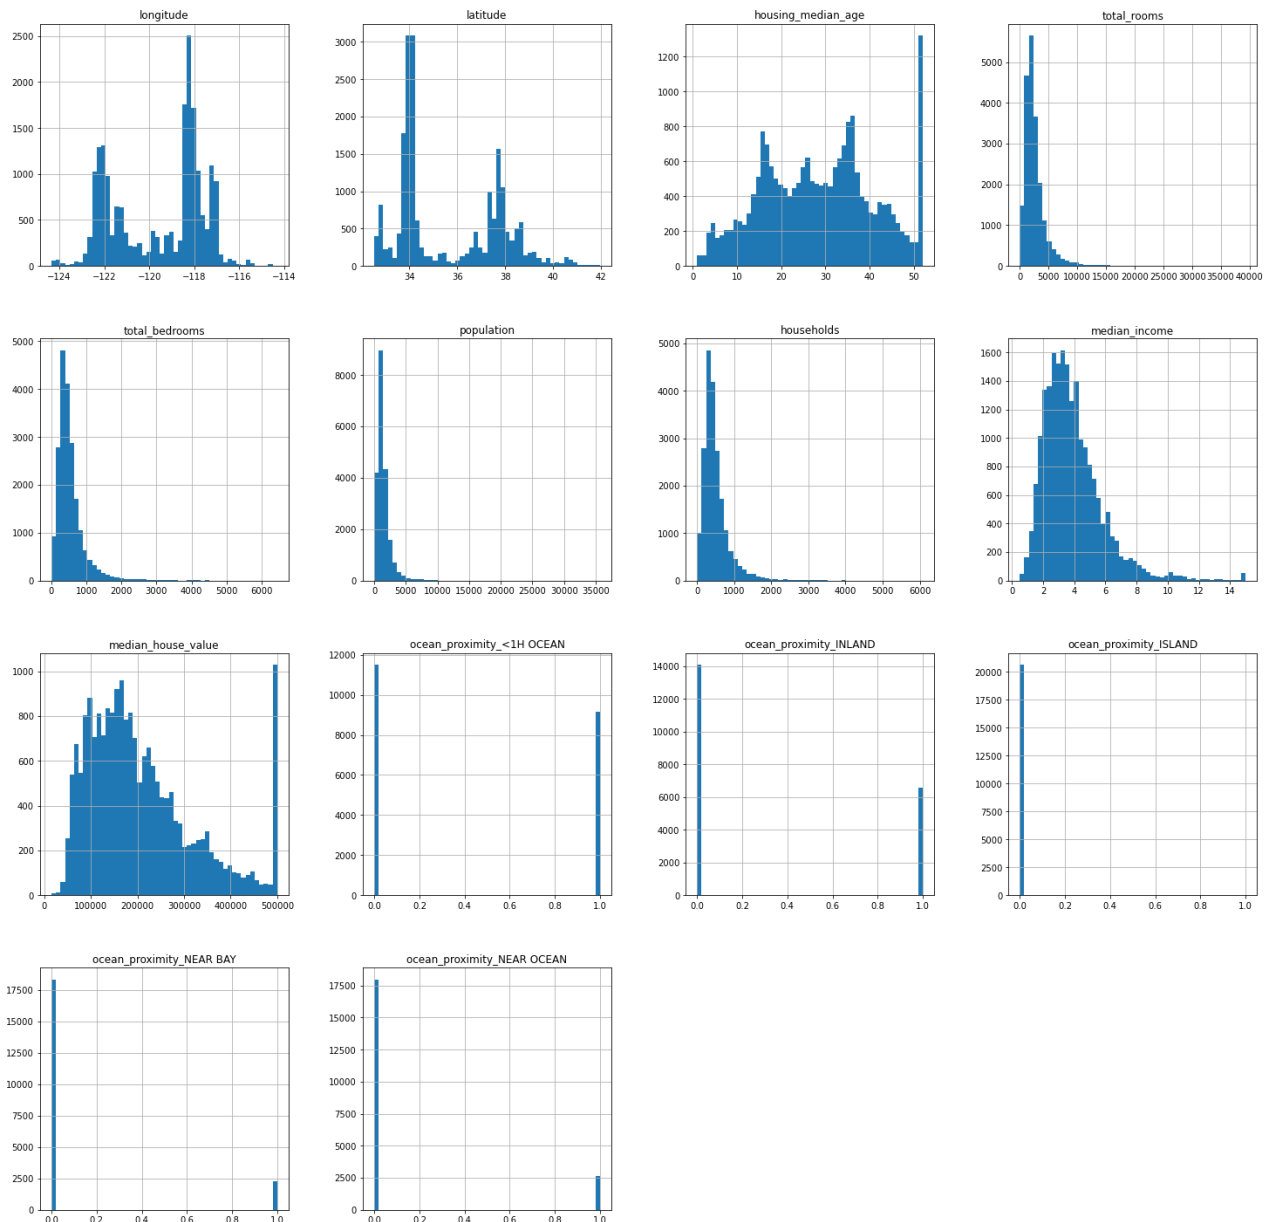
(20640, 13) (20640,)

```

```

In [114... df_housing_final.hist(figsize=(25,25),bins=50);

```



```
In [115... #Split the data into 80% training dataset and 20% test dataset.
x_train,x_test,y_train,y_test=train_test_split(x_input,y_output,test_size=0.2,random_st
print( x_train.shape,x_test.shape, y_train.shape, y_test.shape)
```

```
(16512, 13) (4128, 13) (16512,) (4128,)
```

```
In [116... #Standardize training and test datasets.
scale= StandardScaler()
scaled_train_data = scale.fit_transform(x_train)
scaled_test_data = scale.fit_transform(x_test)
print(scaled_train_data , scaled_test_data)

[[-1.42250942  0.97229046  1.85890297 ... -0.01740407  2.82640555
 -0.38546202]
 [-1.38265919  1.08459626  1.06434823 ... -0.01740407  2.82640555
 -0.38546202]
 [-0.8297373  1.06119922 -1.0014941  ... -0.01740407 -0.35380627
 -0.38546202]
 ...
 [ 0.65468363 -0.79652586  1.06434823 ... -0.01740407 -0.35380627
 -0.38546202]
 [ 1.20262424 -0.89011402 -1.47822694 ... -0.01740407 -0.35380627
 -0.38546202]
```

```

[-1.30794002  1.00972573  0.50815991 ... -0.01740407  2.82640555
 -0.38546202]] [[ 0.59953305 -0.73685251  0.81226638 ...  0.          -0.35109159
 -0.38047173]
[-0.11505424  0.53929953  0.65331708 ...  0.          -0.35109159
 -0.38047173]
[-1.44358273  0.9850144   1.36858896 ...  0.          2.84825961
 -0.38047173]
...
[-1.4184212   0.92402184 -0.22090411 ...  0.          -0.35109159
 2.62831619]
[ 0.73037298 -0.72277731  1.05069034 ...  0.          -0.35109159
 -0.38047173]
[ 1.09269893 -0.76969466  1.84543688 ...  0.          -0.35109159
 -0.38047173]]

```

```

In [117... #Perform Linear Regression on training data.
regressor = LinearRegression()
regressor.fit(x_train, y_train)
print(regressor.intercept_, regressor.coef_)

-2224231.212392007 [-2.65375452e+04 -2.51693181e+04  1.06947068e+03 -5.49854147e+00
 7.81334978e+01 -3.84586986e+01  7.13344968e+01  3.93198087e+04
 -2.41189059e+04 -6.34495807e+04  1.31959049e+05 -2.62977843e+04
 -1.80927777e+04]

```

```

In [118... #Predict output for test dataset using the fitted model.
y_predict = regressor.predict(x_test)

```

```

In [119... #Print root mean squared error (RMSE) from Linear Regression.
np.sqrt(metrics.mean_squared_error(y_test, y_predict))

```

```

Out[119... 68949.62451074323

```

Bonus exercise: Perform Linear Regression with one independent variable

```

In [120... #Extract just the median_income column from the independent variables (from X_train and

x_train_mi = x_train['median_income']
x_train_mi = x_train_mi.values.reshape(-1,1)
x_test_mi = x_test['median_income']
x_test_mi = x_test_mi.values.reshape(-1,1)
y_train = y_train.values.reshape(-1,1)
print(x_train_mi.shape, x_test_mi.shape)

```

```

(16512, 1) (4128, 1)

```

```

In [121... #Perform Linear Regression to predict housing values based on median_income.
regressor.fit(x_train_mi, y_train)
print(regressor.intercept_, regressor.coef_)

```

```

[44721.83362107] [[42055.4573838]]

```

```

In [122... #Predict output for test dataset using the fitted model.

```

```

y_predict_mi = regressor.predict(x_test_mi)
y_predict_mi = y_predict_mi.reshape(-1,1)

```

```

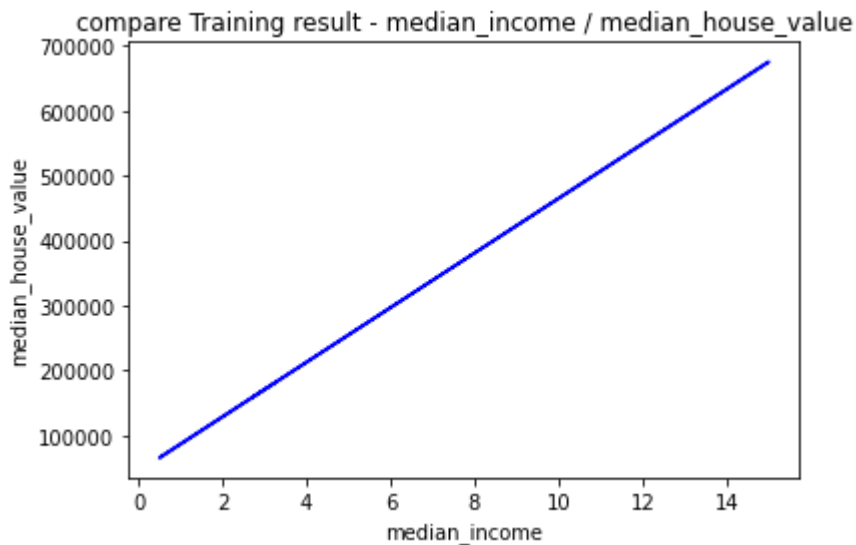
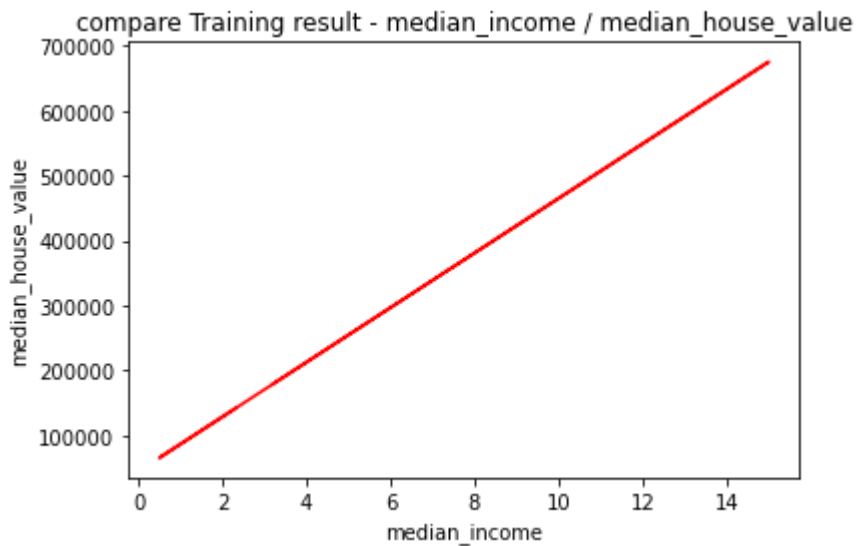
In [123... np.sqrt(metrics.mean_squared_error(y_test, y_predict_mi))

```

Out[123... 83228.17849797675


```
In [124... #Plot the fitted model for training data as well as for test data to check if the fitte
plt.plot (x_train_mi, regressor.predict(x_train_mi), color = 'red')
plt.title ('compare Training result - median_income / median_house_value')
plt.xlabel('median_income')
plt.ylabel('median_house_value')
plt.show()

plt.plot (x_test_mi, regressor.predict(x_test_mi), color = 'blue')
plt.title ('compare Training result - median_income / median_house_value')
plt.xlabel('median_income')
plt.ylabel('median_house_value')
plt.show()
```



```
In [125... df_housing_final.plot(kind='scatter',x='median_income',y='median_house_value')
plt.plot(x_test_mi,y_predict_mi,c='red',linewidth=2)
```

Out[125... [<matplotlib.lines.Line2D at 0x24769cfba30>]



In []: